

# New York Taxi Data Analysis

Section 0501

BUSI758B

Ridhima Amruthesh  
Kiruthika Sankaran  
Abhinav Tummala  
Shreyash Suryawanshi  
Kunal Kamat Tarcar

# Dataset Description

- Agency : Taxi and Limousine Commission (TLC)
- Two types of taxis:
  - Yellow
  - Green

```
[74] yellow_april = spark.read.csv("yellow_tripdata_2018-04.csv", header=True, inferSchema=True)
      yellow_may = spark.read.csv("yellow_tripdata_2018-05.csv", header=True, inferSchema=True)
```

```
print((yellow_combined.count(), len(yellow_combined.columns)))
```

```
(18529578, 17)
```

```
print((green_combined.count(), len(green_combined.columns)))
```

```
(1597317, 19)
```

```
[76] yellow_combined.printSchema()
```

```
↳ root
  |-- VendorID: integer (nullable = true)
  |-- tpep_pickup_datetime: timestamp (nullable = true)
  |-- tpep_dropoff_datetime: timestamp (nullable = true)
  |-- passenger_count: integer (nullable = true)
  |-- trip_distance: double (nullable = true)
  |-- RatecodeID: integer (nullable = true)
  |-- store_and_fwd_flag: string (nullable = true)
  |-- PULocationID: integer (nullable = true)
  |-- DOLocationID: integer (nullable = true)
  |-- payment_type: integer (nullable = true)
  |-- fare_amount: double (nullable = true)
  |-- extra: double (nullable = true)
  |-- mta_tax: double (nullable = true)
  |-- tip_amount: double (nullable = true)
  |-- tolls_amount: double (nullable = true)
  |-- improvement_surcharge: double (nullable = true)
  |-- total_amount: double (nullable = true)
```

# Descriptive Analytics

## Elimination:

- 0.8% data has 0 passengers.
- No Rate Code ID named 99 in Data Dictionary.

```
yellow_combined.groupBy('passenger_count').count().show()
```

passenger_count	count
1	13201488
6	505987
3	756841
5	852545
9	59
4	356618
8	65
7	67
2	2710998
0	144910

```
[ ] yellow_combined.groupBy('RateCodeID').count().show()
```

RateCodeID	count
1	16557633
6	11
3	1280
5	1925
4	1176
2	5
99	7

# Descriptive Analytics

## Outlier Detection:

- Minimum fare amount is 2.5.
- Deleting outliers beyond 1.5 Interquartile range.

```
cols = ['trip_distance', 'fare_amount', 'extra', 'mta_tax', 'tip_amount',
        'tolls_amount', 'improvement_surcharge', 'total_amount']
green_bounds = {}
for col in cols:
    quantiles = green_combined.approxQuantile(
        col, [0.25, 0.75], 0.05
    )
    IQR = quantiles[1] - quantiles[0]
    green_bounds[col] = [
        quantiles[0] - 1.5 * IQR,
        quantiles[1] + 1.5 * IQR
    ]
```

```
[80] bounds
```

```
{'fare_amount': [-5.25, 28.75], 'trip_distance': [-3.2649999999999992, 8.375]}
```

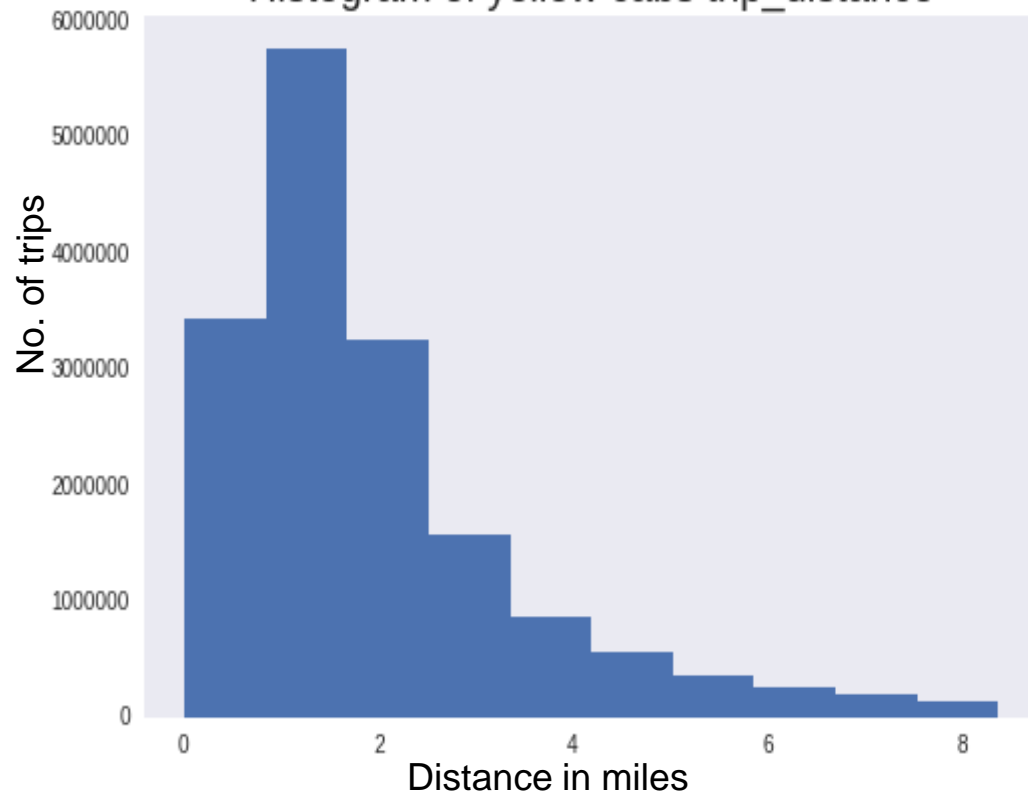
```
+-----+
|      trip_distance      |
+-----+
|          18529578        |
| 2.9743718156991483      |
| 3.843483044976954      |
|              0.0        |
|             943.5        |
+-----+

+-----+
|      fare_amount        |
+-----+
|          18529578        |
| 13.185788835558036      |
| 98.41656868783119      |
|             -485.0       |
|          349026.72       |
+-----+
```

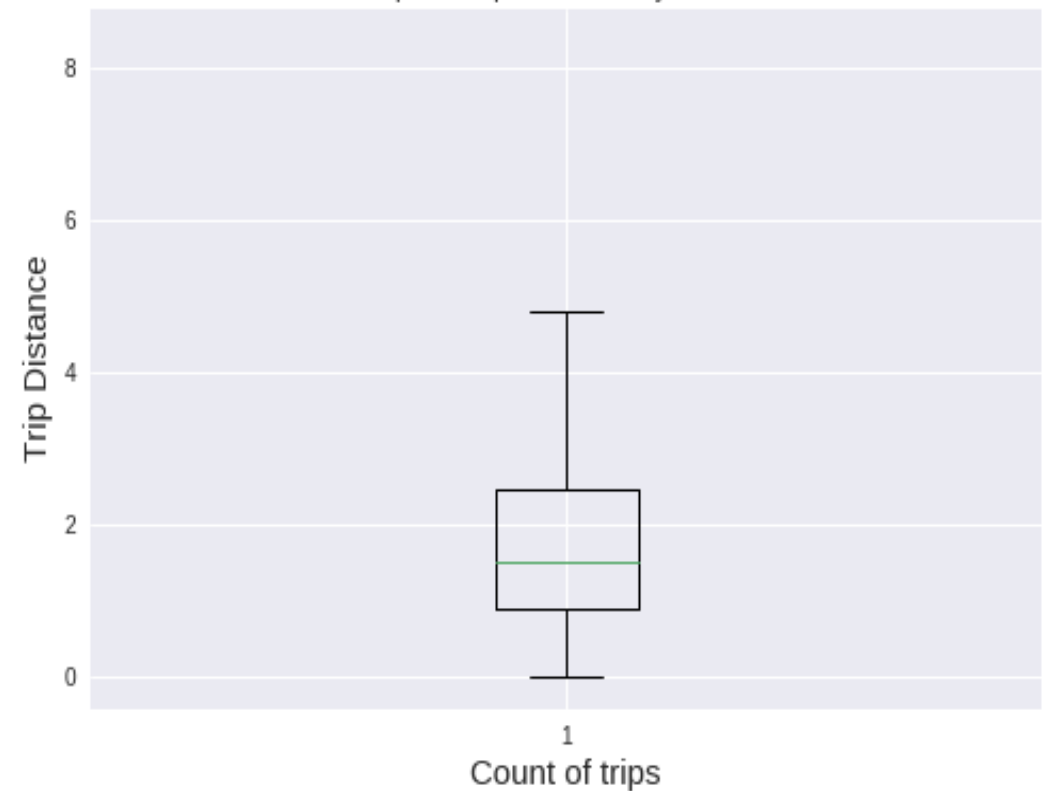
# Descriptive Analytics

## Trip Distance Analysis:

Histogram of yellow cabs trip\_distance



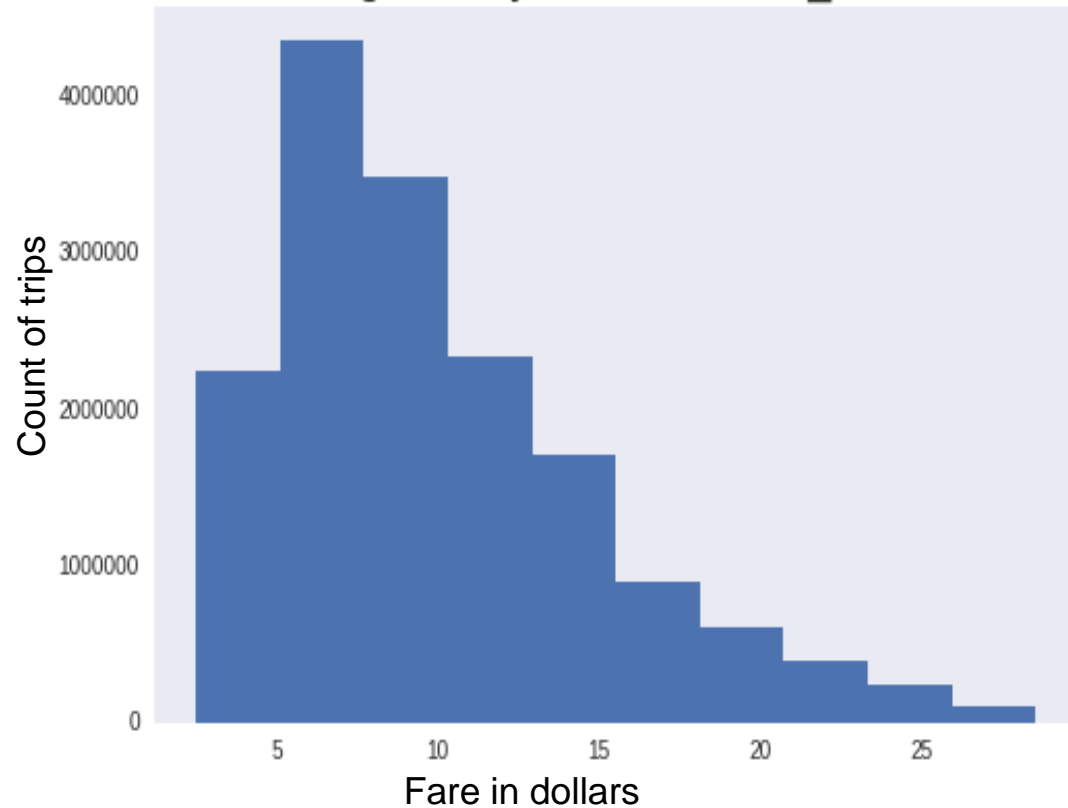
Box plot of trip distance for yellow cabs



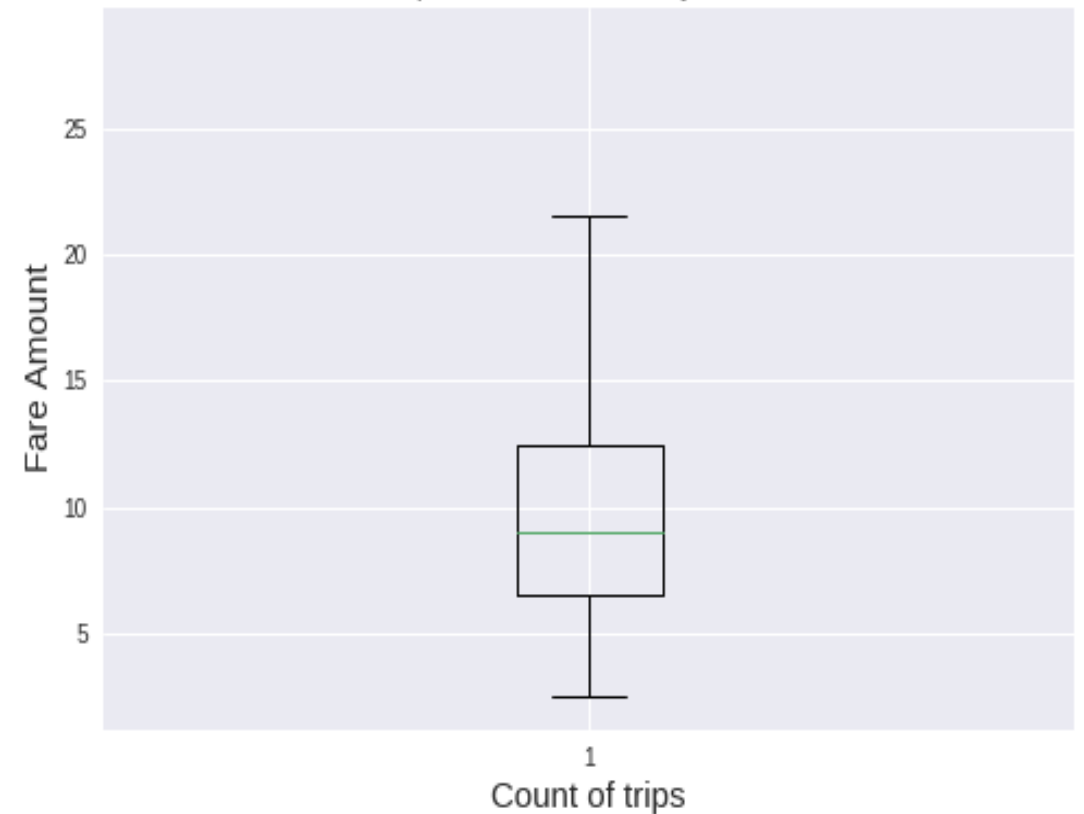
# Descriptive Analytics

## Fare Amount Analysis:

Histogram of yellow cabs fare\_amount



Box plot of fare amount for yellow cabs





# Feature Engineering

## New Features:

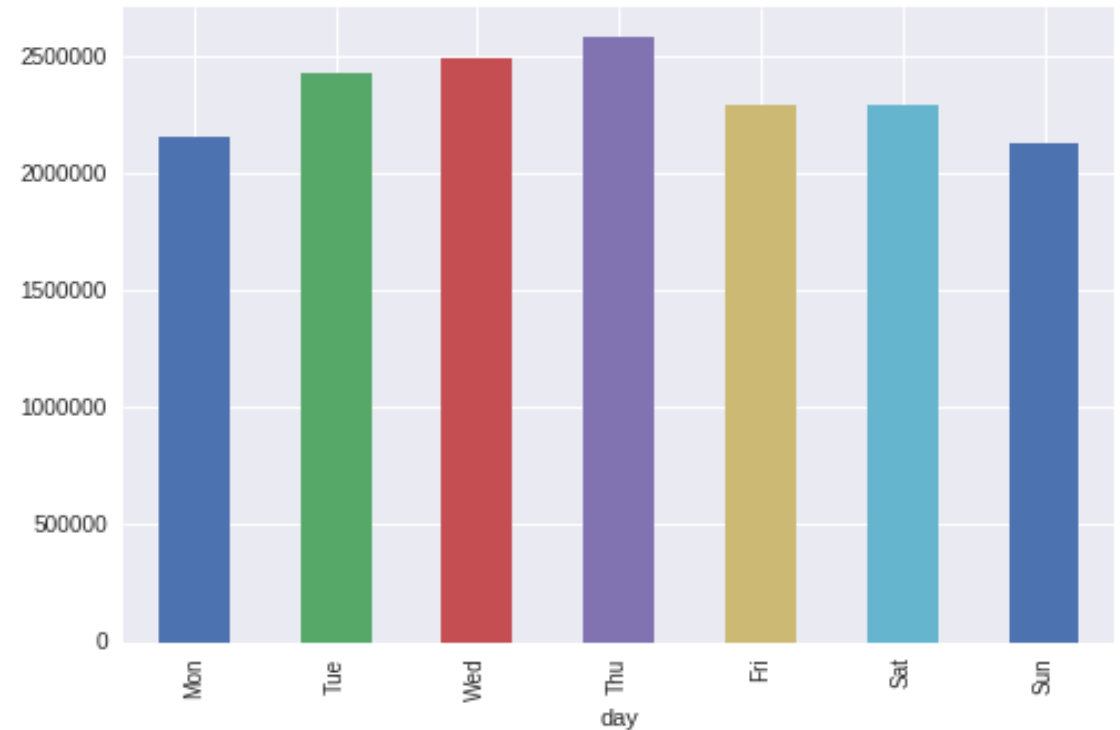
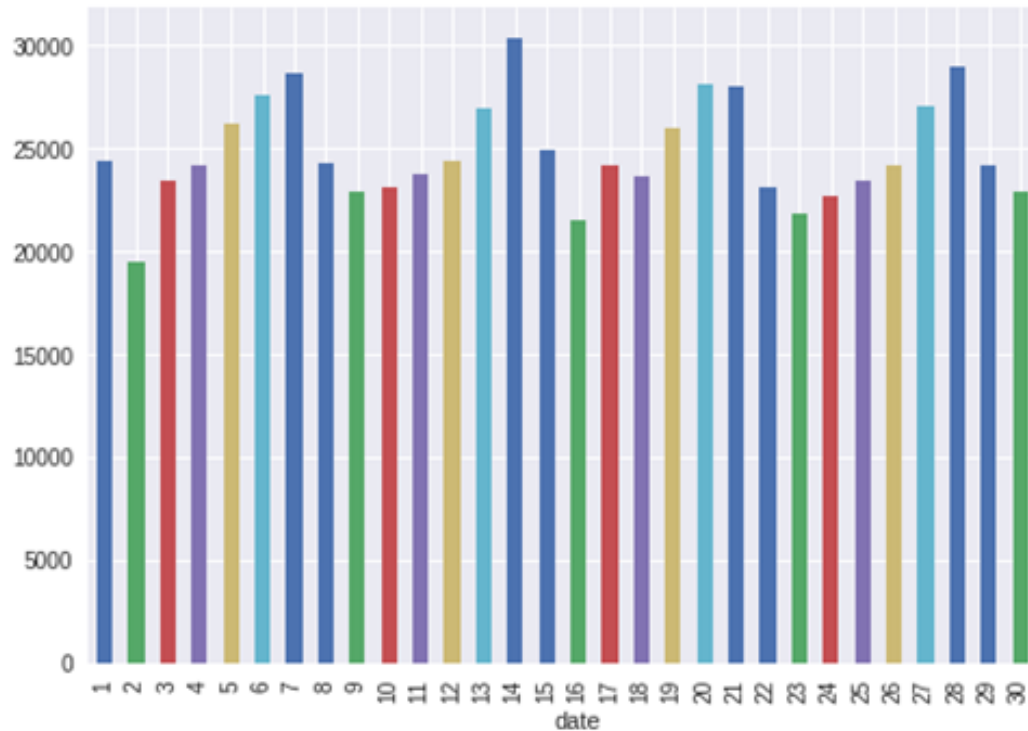
- Time related - hour, date, day, weekend
- Duration of the ride
- Average Taxi speed

tpep_pickup_datetime	hour	date_yellow	day_number	day	is_weekend
2018-04-01 00:42:17	0	1	7	Sun	1
2018-04-01 00:10:47	0	1	7	Sun	1
2018-04-01 00:39:01	0	1	7	Sun	1
2018-04-01 00:44:42	0	1	7	Sun	1
2018-04-01 00:55:11	0	1	7	Sun	1

duration	speed
1.9333333333333333	9.310344827586206
10.266666666666667	13.44155844155844
12.733333333333333	7.539267015706807
3.7333333333333334	7.232142857142857
8.116666666666667	17.445585215605746

# Inferential Analytics

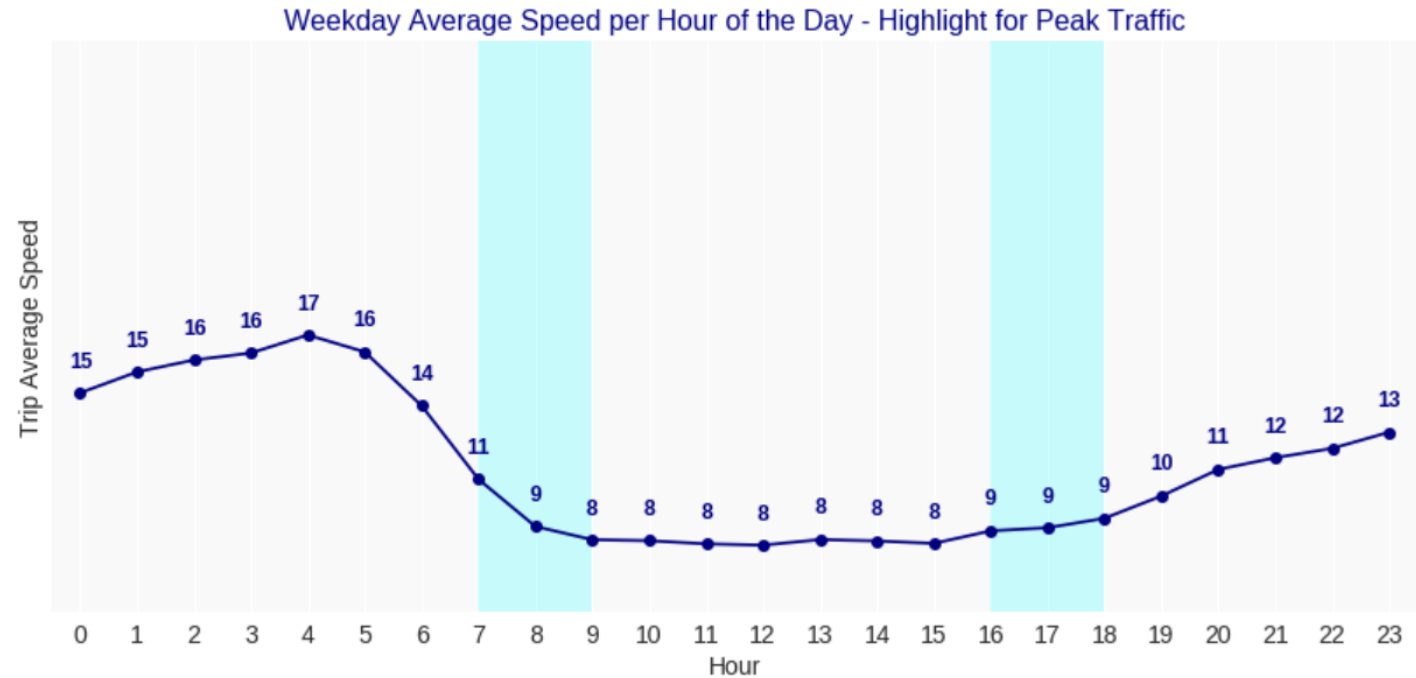
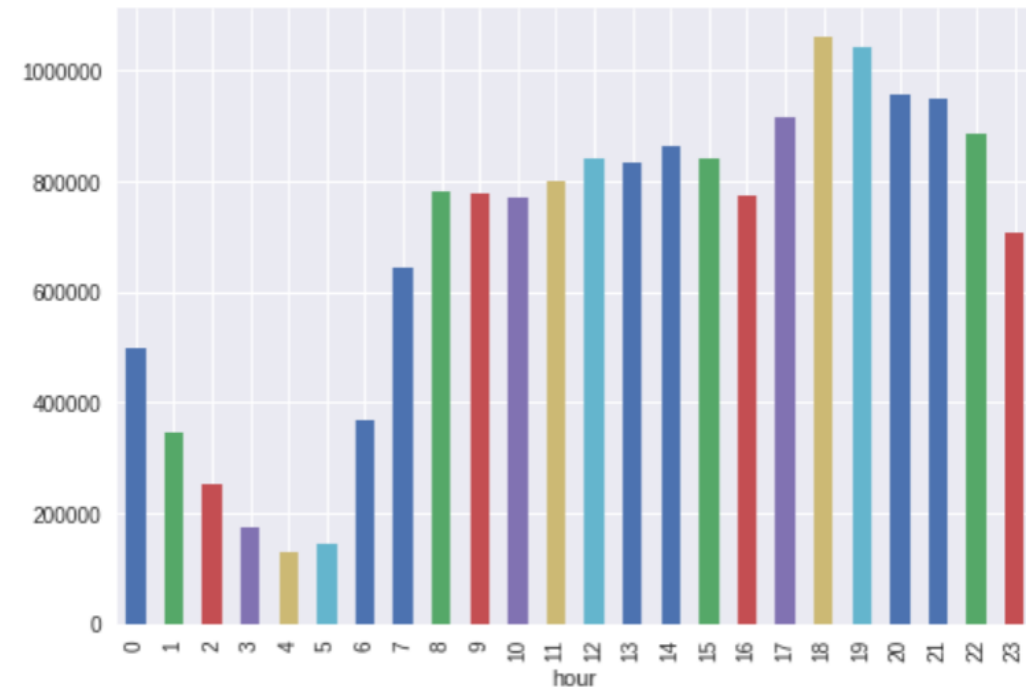
- Weekly trend observed in the daily number of bookings in April.
- Maximum bookings on Thursdays.





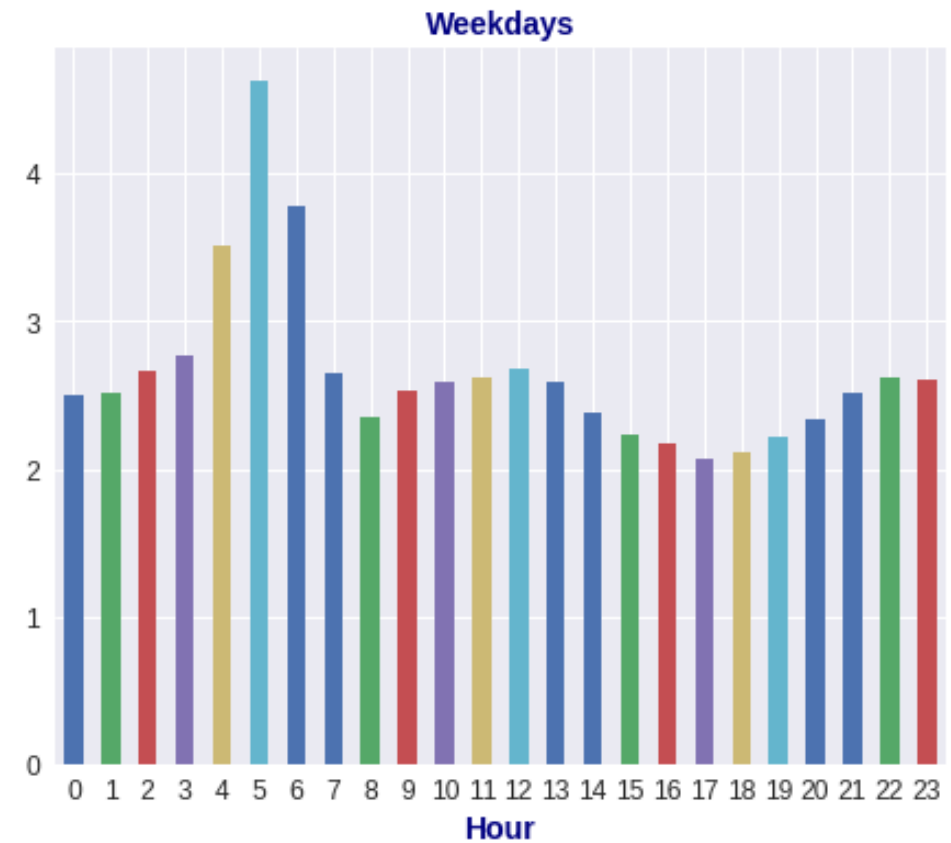
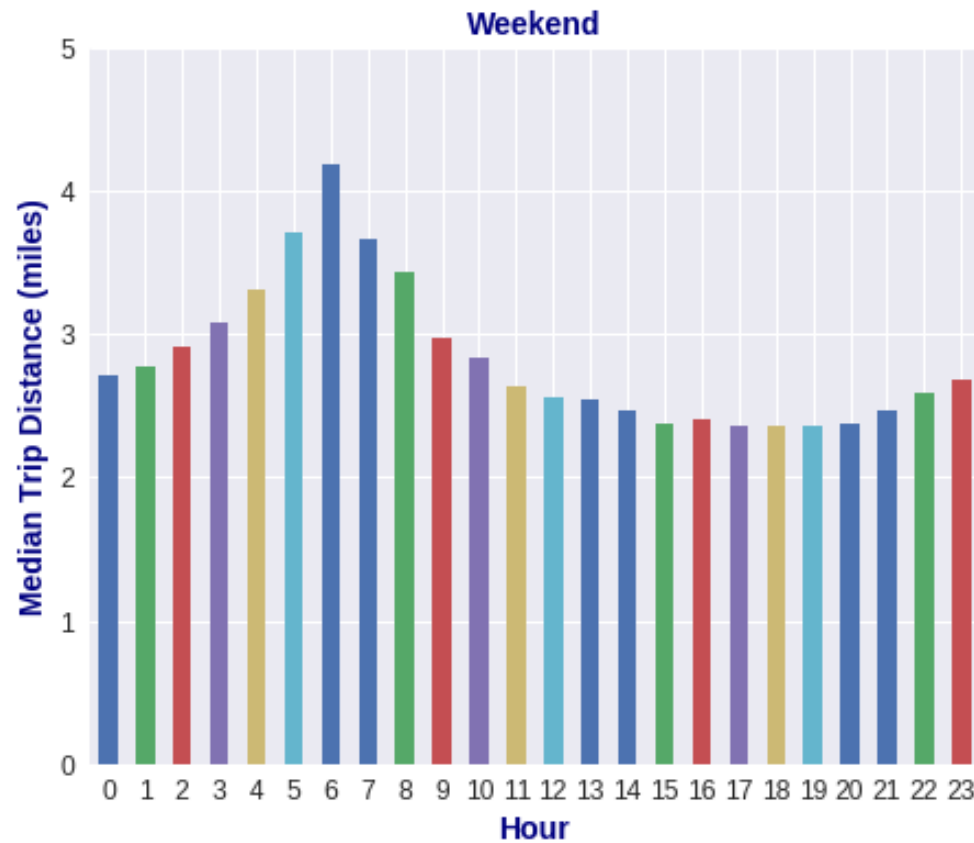
# Inferential Analytics

Demand for yellow cabs per hour of day  
- Weekdays Peak Hours from 5pm-9pm



# Inferential Analytics

Median Trip Distance is more during off-peak hours.



# Inferential Analytics

- Total number of trips in Manhattan is significantly higher than Queens or Brooklyn.
- However, number of disputes is more in Queens and Brooklyn.
- 36.6 % of disputes happen in Queens.

**DISPUTE** →

Payment_type	count
1	739608
3	3106
5	41
4	2407
2	623425

PUBorough	count
Queens	882
Brooklyn	834
Manhattan	558
Bronx	130
Unknown	3



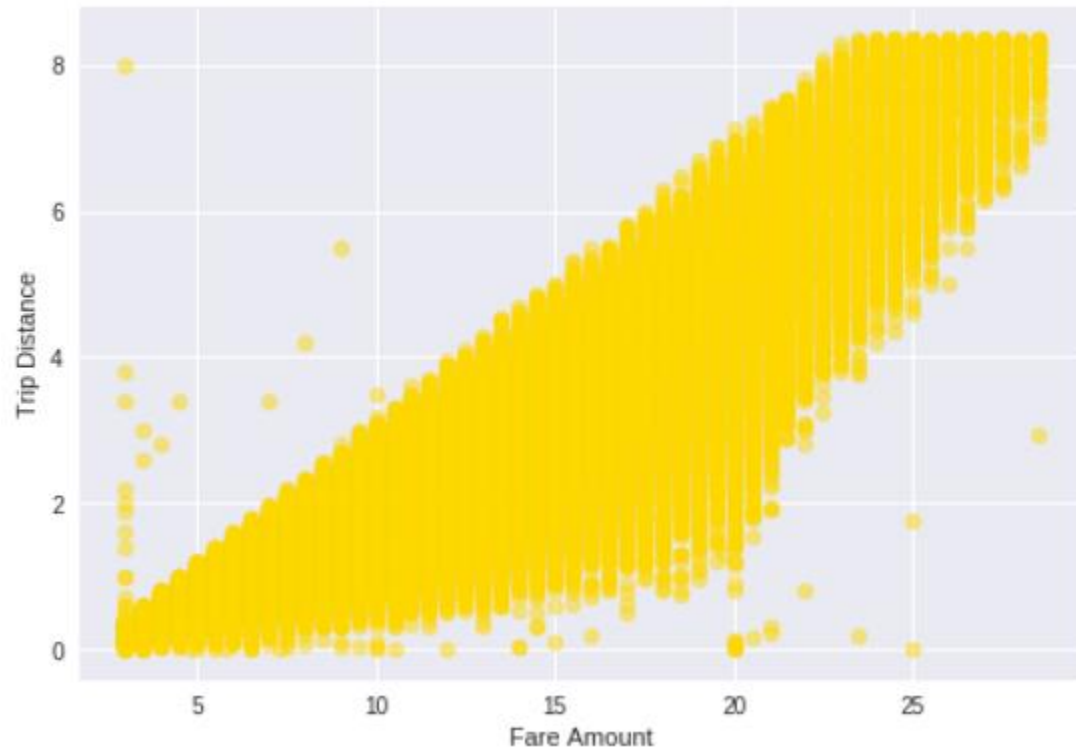
+-----+	
PUBorough   count-Weekend	
+-----+	
Brooklyn	138815
Manhattan	124749
Queens	124502
Bronx	16748
Staten Island	23
+-----+	

+-----+	
PUBorough   count-Weekday	
+-----+	
Manhattan	346416
Brooklyn	284177
Queens	282237
Bronx	50096
Staten Island	53
+-----+	

+-----+	
PUZone   count-Weekend	
+-----+	
East Village	182089
Upper East Side S...	152946
Clinton East	149356
Penn Station/Madi...	145279
Lincoln Square East	142553
+-----+	
only showing top 5 rows	

+-----+	
PUZone   count-Weekday	
+-----+	
Upper East Side S...	599121
Upper East Side N...	524910
Midtown Center	517182
Midtown East	467894
Murray Hill	413015
+-----+	
only showing top 5 rows	

# Fare Amount VS Trip Distance, Duration



```
[ ] yellow_combined.corr('trip_distance', 'fare_amount')
```

```
0.9171635404358489
```

```
yellow_combined.corr('duration', 'fare_amount')
```

```
0.9198622648136741
```







# Recommendation

- Taxi fares are fixed throughout the year, regardless of seasonality.
- Competitors are capitalizing on peak hours when demand is greater.
- Variable Pricing - Capitalize on peak hours.
- Location Targeting based on weekday/weekend and hour of the day.



THANK YOU

# Appendix - Code

```
!apt-get install openjdk-8-jdk-headless -qq > /dev/null
!wget -q https://www-eu.apache.org/dist/spark/spark-2.4.0/spark-2.4.0-bin-hadoop2.7.tgz
!tar xf spark-2.4.0-bin-hadoop2.7.tgz
!pip install -q findspark

# Set environmental variables
import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
#os.environ["SPARK_HOME"] = "/content/spark-2.2.1-bin-hadoop2.7"
os.environ["SPARK_HOME"] = "/content/spark-2.4.0-bin-hadoop2.7"
```

```
import findspark
findspark.init()
from pyspark.sql import SparkSession

spark = SparkSession.builder.master("local[*]").getOrCreate()

spark

from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('NYC').getOrCreate()
```

# Appendix - Code

```
!wget https://s3.amazonaws.com/nyc-
tlc/trip+data/yellow_tripdata_2018-04.csv
!wget https://s3.amazonaws.com/nyc-
tlc/trip+data/yellow_tripdata_2018-05.csv
print(spark.version)
yellow_april = spark.read.csv("yellow_tripdata_2018-04.csv",
header=True, inferSchema=True)
yellow_may = spark.read.csv("yellow_tripdata_2018-05.csv",
header=True, inferSchema=True)
green_april = spark.read.csv("green_tripdata_2018-04.csv",
header=True, inferSchema=True)
green_may = spark.read.csv("green_tripdata_2018-05.csv",
header=True, inferSchema=True)
yellow_combined = yellow_april.union(yellow_may)
yellow_combined.printSchema()
green_combined = green_april.union(green_may)
yellow_combined =
yellow_combined.filter(yellow_combined["VendorID"]!=4)
```

```
yellow_combined =
yellow_combined.filter(yellow_combined["passenger_count"]!=0)
green_combined =
green_combined.filter(green_combined["passenger_count"]!=0)
#Outliers Detection
cols = ['trip_distance', 'fare_amount', 'extra', 'mta_tax', 'tip_amount',
        'tolls_amount', 'improvement_surcharge', 'total_amount']
bounds = {}
for col in cols:
    quantiles = yellow_combined.approxQuantile(
        col, [0.25, 0.75], 0.05
    )
    IQR = quantiles[1] - quantiles[0]
    bounds[col] = [
        quantiles[0] - 1.5 * IQR,
        quantiles[1] + 1.5 * IQR
    ]
bounds
```

# Appendix - Code

```
yellow_combined =
yellow_combined.filter((yellow_combined["trip_distance"] > 0) &
                        (yellow_combined["trip_distance"] <
bounds['trip_distance'][1]))
yellow_combined =
yellow_combined.filter((yellow_combined["fare_amount"] > 2.5) &
                        (yellow_combined["fare_amount"] <
bounds['fare_amount'][1]))
print((yellow_combined.count(), len(yellow_combined.columns)))
cols = ['trip_distance', 'fare_amount', 'extra', 'mta_tax', 'tip_amount',
        'tolls_amount', 'improvement_surcharge', 'total_amount']
green_bounds = {}
for col in cols:
    quantiles = green_combined.approxQuantile(
        col, [0.25, 0.75], 0.05
    )
    IQR = quantiles[1] - quantiles[0]
    green_bounds[col] = [
        quantiles[0] - 1.5 * IQR,
        quantiles[1] + 1.5 * IQR
```

```
green_combined =
green_combined.filter((green_combined["trip_distance"] > 0) &
                        (green_combined["trip_distance"] <
green_bounds['trip_distance'][1]))
green_combined =
green_combined.filter((green_combined["fare_amount"] > 2.5) &
                        (green_combined["fare_amount"] <
green_bounds['fare_amount'][1]))
yellow_combined =
yellow_combined.filter(yellow_combined["RateCodeID"] != 99)
from pyspark.sql import functions as F
timeFmt = "yyyy-MM-dd'T'HH:mm:ss.SSS"
timeDiff = (F.unix_timestamp('tpep_dropoff_datetime',
format=timeFmt)
            - F.unix_timestamp('tpep_pickup_datetime', format=timeFmt))
yellow_combined = yellow_combined.withColumn("duration",
timeDiff/60)
```

# Appendix - Code

```
cols = ['duration']
time_bounds = {}
for col in cols:
    quantiles = yellow_combined.approxQuantile(
        col, [0.25, 0.75], 0.05
    )
    IQR = quantiles[1] - quantiles[0]
    time_bounds[col] = [
        quantiles[0] - 1.5 * IQR,
        quantiles[1] + 1.5 * IQR
    ]
time_bounds
yellow_combined =
yellow_combined.filter((yellow_combined["duration"] > 0) &
(yellow_combined["duration"] < time_bounds['duration'][1]))
from pyspark.sql import functions as F
timeFmt = "yyyy-MM-dd'T'HH:mm:ss.SSS"
```

```
timeDiff = (F.unix_timestamp('lpep_dropoff_datetime',
format=timeFmt)
            - F.unix_timestamp('lpep_pickup_datetime', format=timeFmt))
green_combined = green_combined.withColumn("duration",
timeDiff/60)
cols = ['duration']
time_bounds = {}
for col in cols:
    quantiles = green_combined.approxQuantile(
        col, [0.25, 0.75], 0.05
    )
    IQR = quantiles[1] - quantiles[0]
    time_bounds[col] = [
        quantiles[0] - 1.5 * IQR,
        quantiles[1] + 1.5 * IQR
    ]
time_bounds
green_combined =
green_combined.filter((green_combined["duration"] > 0) &
```



# Appendix - Code

```
from pyspark.sql.functions import date_format
from pyspark.sql.functions import isnan, when, count, col
```

```
date_green = green_combined.select('lpep_pickup_datetime',
                                   date_format('lpep_pickup_datetime', 'H').alias('hour'),
                                   date_format('lpep_pickup_datetime', 'd').alias('date'),
                                   date_format('lpep_pickup_datetime',
                                   'u').alias('day_number'), date_format('lpep_pickup_datetime',
                                   'E').alias('day'))
#green.show()

date_green = date_green.withColumn('is_weekend',
when((date_green.day == 'Sun') | (date_green.day ==
'Sat'),1).otherwise(0))

# Convert string to numeric
```

```
from pyspark.sql.types import IntegerType

date_green = date_green.withColumn("hour",
date_green["hour"].cast(IntegerType()))

date_green = date_green.withColumn("date",
date_green["date"].cast(IntegerType()))
```

```
date_green = date_green.withColumn("day_number",
date_green["day_number"].cast(IntegerType()))
```

```
date_green = date_green.withColumn("is_weekend",
date_green["is_weekend"].cast(IntegerType()))
```

```
from pyspark.sql.functions import date_format
from pyspark.sql.functions import isnan, when, count, col
```

```
date_yellow = yellow_combined.select('tpep_pickup_datetime',
                                     date_format('tpep_pickup_datetime',
                                     'H').alias('hour'),
                                     date_format('tpep_pickup_datetime',
                                     'd').alias('date_yellow'),
                                     date_format('tpep_pickup_datetime',
                                     'u').alias('day_number'), date_format('tpep_pickup_datetime',
                                     'E').alias('day'))
```

```
#yellow_combined.show()
```

```
date_yellow =
date_yellow.withColumn('is_weekend',when((date_yellow.day ==
'Sun') | (date_yellow.day == 'Sat'),1).otherwise(0))
```

```
# Convert string to numeric
```

# Appendix - Code

```
from pyspark.sql.types import IntegerType

date_yellow = date_yellow.withColumn("hour",
date_yellow["hour"].cast(IntegerType()))

date_yellow = date_yellow.withColumn("date_yellow",
date_yellow["date_yellow"].cast(IntegerType()))

date_yellow = date_yellow.withColumn("day_number",
date_yellow["day_number"].cast(IntegerType()))

date_yellow = date_yellow.withColumn("is_weekend",
date_yellow["is_weekend"].cast(IntegerType()))

from pyspark.sql.functions import monotonically_increasing_id

green_combined = green_combined.drop("lpep_pickup_datetime")
green_combined = green_combined.withColumn("id",
monotonically_increasing_id())

date_green = date_green.withColumn("id",
monotonically_increasing_id())

green_combined = date_green.join(green_combined, "id",
"outer").drop("id")
```

```
green_combined.show()
```

```
from pyspark.sql.functions import monotonically_increasing_id

yellow_combined = yellow_combined.drop("tpep_pickup_datetime")
yellow_combined = yellow_combined.withColumn("id",
monotonically_increasing_id())

date_yellow = date_yellow.withColumn("id",
monotonically_increasing_id())

yellow_combined = date_yellow.join(yellow_combined, "id",
"outer").drop("id")

yellow_combined.show(5)

yellow_combined = yellow_combined.withColumn("speed",
yellow_combined['trip_distance']*60/yellow_combined['duration'])

green_combined = green_combined.withColumn("speed",
green_combined['trip_distance']*60/green_combined['duration'])
```

# Appendix - Code

## #Visualizations

### ##Histogram

```
import matplotlib.pyplot as plt
```

```
# Show histogram of the 'C1' column
```

```
bins, counts = yellow.select('trip_distance').rdd.flatMap(lambda x:  
x).histogram(10)
```

```
plt.hist(bins[:-1], bins=bins, weights=counts)
```

```
bins, counts = green_new.select('trip_distance').rdd.flatMap(lambda x:  
x).histogram(10)
```

```
plt.hist(bins[:-1], bins=bins, weights=counts)
```

### ##Barplot

```
# Create dataframe with frequencies with date
```

```
hour_plot = joined1.groupBy('hour').count()
```

```
# Convert the dataframe to pandas dataframe
```

```
hour_plot = hour_plot.toPandas()
```

```
# Set the hour column as index
```

```
hour_plot = hour_plot.set_index('hour')
```

```
hour_plot.sort_index(inplace=True)
```

```
hour_plot.T.squeeze().plot.bar()
```

### ##Boxplot

```
# BoxPlot
```

```
d2 = yellow.select('fare_amount').rdd.flatMap(lambda x: x).collect()
```

```
fig1, ax1 = plt.subplots()
```

```
ax1.set_title('Basic Plot')
```

```
ax1.boxplot(d2)
```

### ##Scatterplot

```
scar = yellow.select('trip_distance','fare_amount')
```

```
scar.show(5)
```

```
s1 = scar.select('fare_amount').rdd.flatMap(lambda x: x).collect()
```

```
s2 = scar.select('trip_distance').rdd.flatMap(lambda x: x).collect()
```

```
plt.scatter(s1,s2, alpha = 0.5, color = 'red', cmap='viridis',  
marker=r'$\clubsuit$', label="Luck")
```

```
plt.xlabel('Fare Amount')
```

```
plt.ylabel('Trip Distance')
```

```
plt.legend(loc='upper left')
```

```
plt.show()
```

# Appendix - Code

#Model

```
df.groupBy("ID").pivot("Text").agg(F.lit(1)).na.fill(0).show()

yellow =
yellow_combined.select('trip_distance','duration','fare_amount','day','
hour')

from pyspark.sql.functions import monotonically_increasing_id
yellow = yellow.withColumn("id", monotonically_increasing_id())
```

```
from pyspark.sql import functions as F
new_yellow = yellow.groupBy("id").pivot("day").agg(F.lit(1)).na.fill(0)
from pyspark.sql import functions as F
new_yellow2 = yellow.groupBy("id").pivot("hour").agg(F.lit(1)).na.fill(0)
dummy_encoded_df = yellow.join(new_yellow, "id", "outer")
dummy_encoded_df = dummy_encoded_df.join(new_yellow2, "id",
"outer").drop("id")
dummy_encoded_df.show(5)
from pyspark.ml.feature import VectorAssembler
```

```
vectorAssembler = VectorAssembler(inputCols =
['trip_distance','duration',
'Mon','Wed','Thu','Fri','Sat','Sun','0','1','2','3','4','6','7','8','9','10','11','12',
'13','14','15','16','17','18','19','20','21','22','23'], outputCol = 'features')
regression_df = vectorAssembler.transform(dummy_encoded_df)
regression_df = regression_df.select(['features', 'fare_amount'])
regression_df.show(3)
```

# Appendix - Code

```
splits = regression_df.randomSplit([0.7, 0.3])
train_df = splits[0]
test_df = splits[1]

from pyspark.ml.regression import LinearRegression

lr = LinearRegression(featuresCol='features', labelCol='fare_amount',
maxIter=10, regParam=0.3, elasticNetParam=0.8)

lr_model = lr.fit(train_df)

print("Coefficients: " + str(lr_model.coefficients))
print("Intercept: " + str(lr_model.intercept))

lr_predictions = lr_model.transform(test_df)

lr_predictions.select("prediction", "fare_amount", "features").show(5)

from pyspark.ml.evaluation import RegressionEvaluator

lr_evaluator = RegressionEvaluator(predictionCol="prediction", \
labelCol="fare_amount", metricName="r2")

print("R Squared (R2) on test data = %g" %
lr_evaluator.evaluate(lr_predictions))
```